

# Data Mining & Machine Learning

CS37300

Purdue University

Sep 8, 2023

# Your First Classifier!

- Let's consider one of the simplest classifiers out there.
- Assume we have a training set  $(x_1, y_1) \dots (x_n, y_n)$
- Now we get a new instance  $x_{\text{new}}$  , **how can we classify it?**
  - Example: Can you recommend a movie, based on user's movie reviews?

# Your First Classifier!

- Let's consider one of the simplest classifiers out there.
- Assume we have a training set  $(x_1, y_1) \dots (x_n, y_n)$
- Now we get a new instance  $x_{\text{new}}$  , **how can we classify it?**
  - Example: Can you recommend a movie, based on user's movie reviews?
- **Simple Solution:**
  - Find the most similar example  $(x, y)$  in the training data and predict the same
    - If you liked "*Fast and Furious*" you'll like "*2 fast 2 furious*"
- Only a single decision is needed: distance metric to compute similarity

# Your First Classifier!

- Let's consider one of the simplest classifiers out there.
- Assume we have a training set  $(x_1, y_1) \dots (x_n, y_n)$
- Now we get a new instance  $x_{\text{new}}$  , **how can we classify it?**
  - Example: Can you recommend a movie, based on user's movie reviews?
- **Simple Solution:**
  - Find the most similar example  $(x, y)$  in the training data and predict the same
    - If you liked "*Fast and Furious*" you'll like "*2 fast 2 furious*"
- Only a single decision is needed: distance metric to compute similarity

$$d(x_1, x_2) = 1 - \frac{x_1 \cap x_2}{x_1 \cup x_2}$$

$$d(x_1, x_2) = \sqrt[2]{(x_1 - x_2)^2}$$

# K Nearest Neighbors

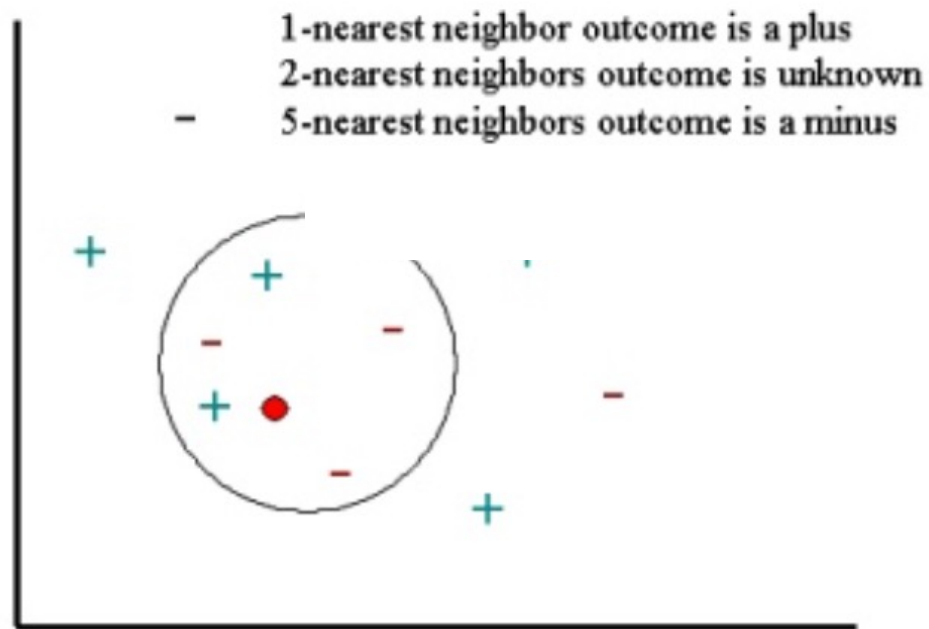
- Can you think about a better way?

# K Nearest Neighbors

- Can you think about a better way?
- We can make the decision by looking at several near examples, not just one.  
*Why would it be better?*

# K Nearest Neighbors

- Can you think about a better way?
- We can make the decision by looking at several near examples, not just one.  
*Why would it be better?*



# K Nearest Neighbors

- **Learning:** just storing the training examples
- **Prediction:**
  - Find the K training example closest to  $\mathbf{x}$
- **Predict a label:**
  - Classification: majority vote
  - Regression: mean value



# K Nearest Neighbors

- **Learning:** just storing the training examples
- **Prediction:**
  - Find the K training example closest to  $\mathbf{x}$
- **Predict a label:**
  - Classification: majority vote
  - Regression: mean value
- KNN is a type of *instance based learning*
- This is called *lazy* learning, since most of the computation is done at prediction time

# Let's analyze KNN

- ***What are the advantages and disadvantages of KNN?***
  - *What should we care about when answering this question?*
- **Complexity**
  - **Space** (*how memory efficient is the algorithm?*)
    - *Why should we care?*
  - **Time** (*computational complexity*)
    - *Both at training time and at test (prediction) time*
- **Expressivity**
  - *What kind of functions can we learn?*

# Let's analyze KNN

- **What are the advantages and disadvantages of KNN?**

- *What should we care about when answering this question?*

- **Complexity**

- **Space** (how memory efficient is the algorithm?)

- *Why should we care?*

- **Time** (computational complexity)

- *Both at training time and at test (prediction)*
      - Training is very fast! But *prediction is slow*
      - $O(dN)$  for  $N$  examples with  $d$  attributes
      - *increases with the number of examples!*

- **Expressivity**

- *What kind of functions can we learn?*

# Let's analyze KNN

- **What are the advantages and disadvantages of KNN?**

- *What should we care about when answering this question?*

- **Complexity**

- **Space** (how memory efficient is the algorithm?)

- *Why should we care?*

KNN needs to maintain all training examples!  
- Datasets can be HUGE

- **Time** (computational complexity)

- *Both at training time and at test (prediction)*
      - Training is very fast! But *prediction is slow*
      - $O(dN)$  for  $N$  examples with  $d$  attributes
      - *increases* with the number of examples!

- **Expressivity**

- *What kind of functions can we learn?*

# Analyzing K Nearest Neighbors


- We discussed the importance of finding a good model space
  - Expressive (we can represent the right model)
  - Constrained (we can search effectively, using the data we have)
- Let's try to characterize the model space, by looking at the **decision boundary**

# Analyzing K Nearest Neighbors

- We discussed the importance of finding a good model space
  - Expressive (we can represent the right model)
  - Constrained (we can search effectively, using the data we have)
- Let's try to characterize the model space, by looking at the **decision boundary**
- **How would it look if  $K=1$ ?**


# Analyzing K Nearest Neighbors

- We discussed the importance of finding a good model space
  - Expressive (we can represent the right model)
  - Constrained (we can search effectively, using the data we have)
- Let's try to characterize the model space, by looking at the **decision boundary**
- **How would it look if  $K=1$ ?**



# Analyzing K Nearest Neighbors

- We discussed the importance of finding a good model space
  - Expressive (we can represent the right model)
  - Constrained (we can search effectively, using the data we have)
- Let's try to characterize the model space, by looking at the **decision boundary**
- **How would it look if  $K=1$ ?**





If we define the model space to be our choice of  $K$   
Does the complexity of the model space increase or decrease with  $K$ ?





# Analyzing K Nearest Neighbors

- Which model has a higher K value?





# Analyzing K Nearest Neighbors

- Which model has a higher K value?
- Which model is more complex?



# Analyzing K Nearest Neighbors

- Which model has a higher K value?
- Which model is more complex?
- Which model is more sensitive to noise?



# Questions

- We know higher  $K$  values result in a smoother decision boundary.
  - Less "jagged" decision regions
  - Total number of regions will be smaller

# Questions

- We know higher  $K$  values result in a smoother decision boundary.
  - Less "jagged" decision regions
  - Total number of regions will be smaller

What will happen if we keep increasing  $K$ , up to the point that  $K=n$  ?  
 $n$  = is the number of examples we have

# How should we determine the value of $K$ ?

- Higher  $K$  values result in less complex functions (less expressive)
- Lower  $K$  values are more complex (more expressive)
- **How can we find the right balance between the two?**

# How should we determine the value of K?

- Higher K values result in less complex functions (less expressive)
- Lower K values are more complex (more expressive)
- **How can we find the right balance between the two?**
- Option 1: Find the K that minimizes the training error.
  - Training error: after learning the classifier, what is the number of errors we get on the training data.
  - What will be this value for  $k=1$ ,  $k=n$ ,  $k=n/2$ ?

# How should we determine the value of K?

- Higher K values result in less complex functions (less expressive)
- Lower K values are more complex (more expressive)
- **How can we find the right balance between the two?**
- Option 1: Find the K that minimizes the training error.
  - Training error: after learning the classifier, what is the number of errors we get on the training data.
  - What will be this value for  $k=1$ ,  $k=n$ ,  $k=n/2$ ?

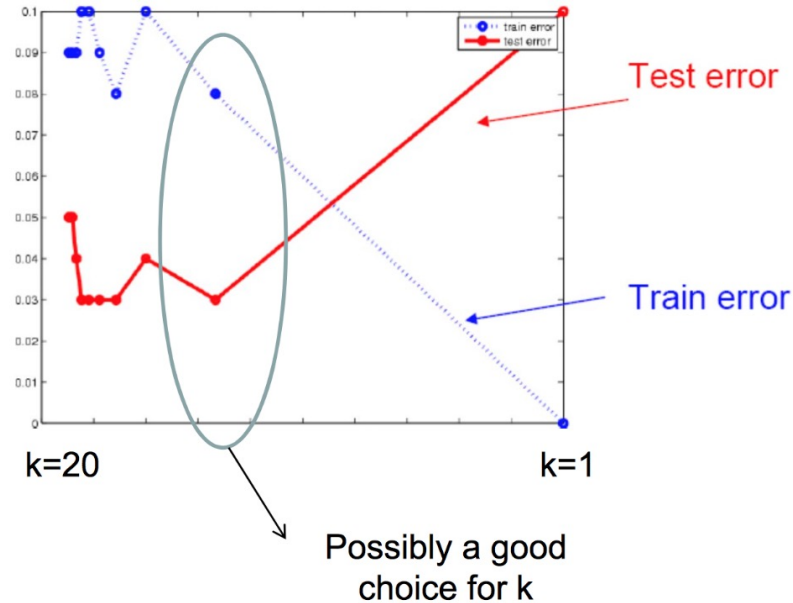
*Is this a good idea?*




# How should we determine the value of K?

- Higher K values result in less complex functions (less expressive)
- Lower K values are more complex (more expressive)
- **How can we find the right balance between the two?**
- Option 1: Find the K that minimizes the training error.
  - Training error: after learning the classifier, what is the number of errors we get on the training data.
  - What will be this value for  $k=1$ ,  $k=n$ ,  $k=n/2$ ? *Is this a good idea?*
- Option 2: Find K that minimizes the **validation error**.
  - Validation error: set aside some of the data (validation set). what is the number of errors we get on the validation data, after training the classifier.

# How should we determine the value of K?



# How should we determine the value of K?



**In general** – using the training error to tune parameters will always result in a more complex hypothesis! **(why?)**

# KNN Practical Consideration

- Finding the right representation is key
  - KNN is very sensitive to irrelevant attributes
- Choosing the right distance metric is important
  - Many options!
  - Popular choices:

– Euclidean distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_2 = \sqrt{\sum_{i=1}^n (\mathbf{x}_{1,i} - \mathbf{x}_{2,i})^2}$$

– Manhattan distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_1 = \sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|$$

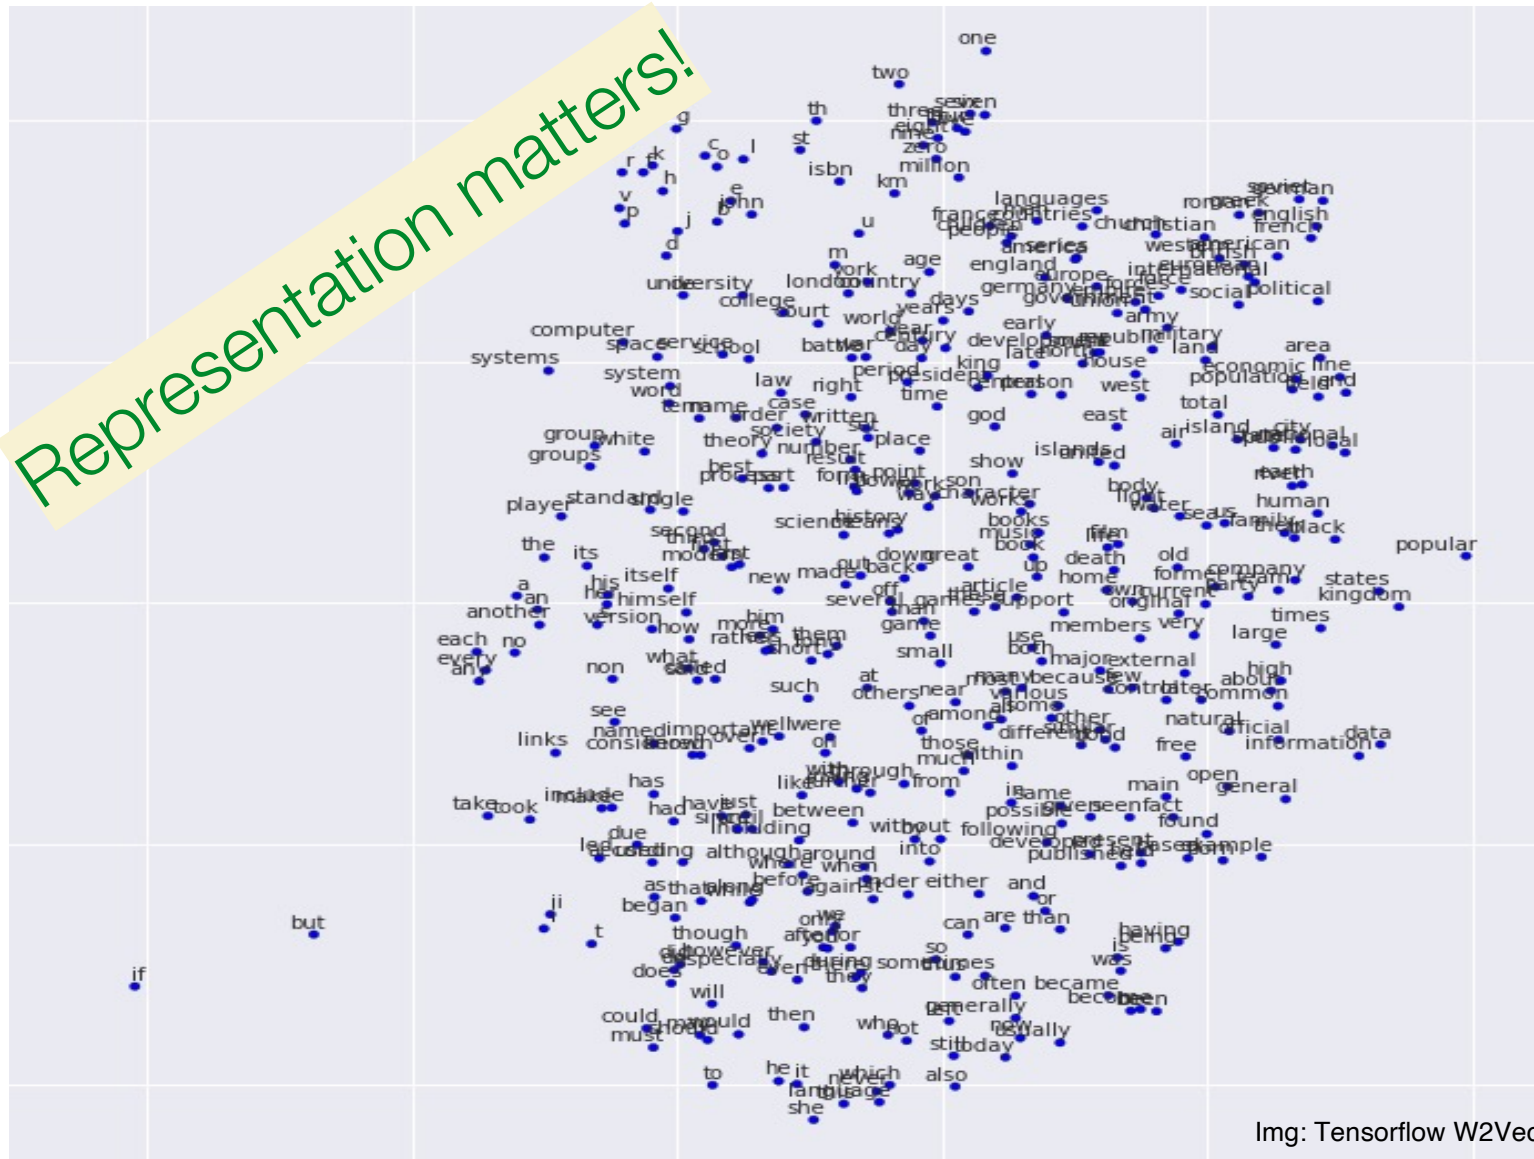
–  $L_p$ -norm

- Euclidean =  $L_2$

- Manhattan =  $L_1$

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_p = \left( \sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|^p \right)^{\frac{1}{p}}$$

Representation matters!



Img: Tensorflow W2Vec


Representation matters!

## France

Word	Cosine distance
spain	0.678515
belgium	0.665923
netherlands	0.652428
italy	0.633130
switzerland	0.622323
luxembourg	0.618033
portugal	0.577154
ruussia	0.571507
germany	0.563291
catalonia	0.534176

# San Francisco

Word	Cosine distance
los_angeles	0.666175
golden_gate	0.571522
oakland	0.557521
california	0.554623
san_diego	0.534939
pasadena	0.519115
seattle	0.512098
taiko	0.507570
houston	0.499762
chicago_illinois	0.491598



Img: Tensorflow W2Vec

# Beyond KNN

- KNN is not a statistical classifier.
- It memorizes the training data, and makes a majority vote over the K closest points.


# Beyond KNN

- KNN is not a statistical classifier.
- It memorizes the training data, and makes a majority vote over the K closest points.
- For example, these two cases are the same:




# Beyond KNN

- KNN is not a statistical classifier.
- It memorizes the training data, and makes a majority vote over the K closest points.
- For example, these two cases are the same:



# Beyond KNN


- KNN is not a statistical classifier.
- It memorizes the training data, and makes a majority vote over the K closest points.
- For example, these two cases are the same:



- What is the difference between the two scenarios?

# Beyond KNN

- KNN is not a statistical classifier.
- It memorizes the training data, and makes a majority vote over the K closest points.
- For example, these two cases are the same:




- What is the difference between the two scenarios?
- How can we reason about it?

# Nearest neighbor

- Strengths:
  - Simple model, easy to implement
- Weaknesses:
  - Inefficient inference: time and space  $O(n)$ 
    - (Inference time improvable with approximations, appropriate data structures)
  - Curse of dimensionality:
    - As number of features increase, you need an exponential increase in the size of the data to ensure that you have "usable" nearby examples for any given data point

# Nearest neighbor

- Strengths:
  - Simple model, easy to implement
- Weaknesses:
  - Inefficient inference: time and space  $O(n)$ 
    - (Inference time improvable with approximations, appropriate data structures)
  - Curse of dimensionality:
    - As number of features increase, you need an exponential increase in the size of the data to ensure that you have "usable" nearby examples for any given data point



# kNN Can Learn ANY Function (with enough data)

- Flexibility: Nearest Neighbor rules can learn **any concept** (with enough data)
- If  $n$  training examples are sampled independently from a distribution,
- if we choose  $k_n \rightarrow \infty$  as  $n \rightarrow \infty$ , but not too fast so  $\frac{k_n}{n} \rightarrow 0$  as  $n \rightarrow \infty$ , then
- kNN's classifier will converge to an optimal predictor
- This is called “universal consistency”