

Data Mining & Machine Learning

CS37300

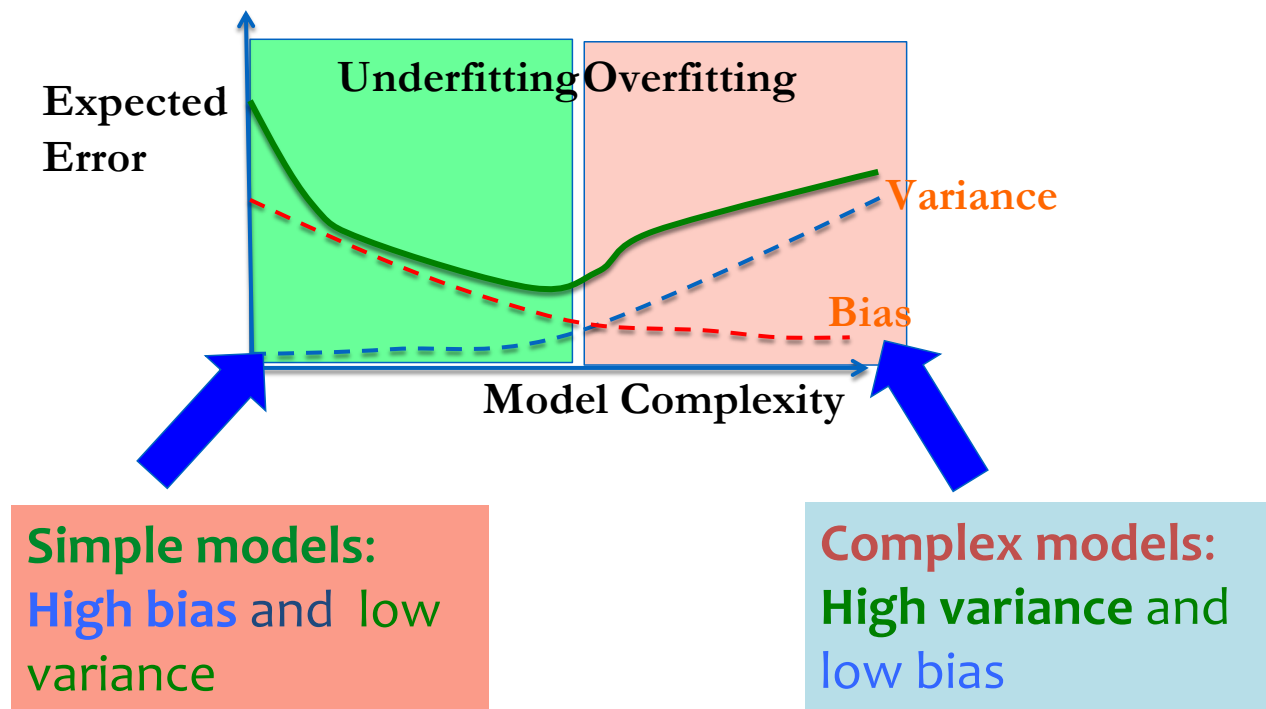
Purdue University

Oct 4, 2023

Bias/Variance (contd)

- Bias:
 - Trust the data less
 - Simpler models
 - Tendency to underfit
 - Could mean need more features
 - “Stable” models
- Variance
 - Trust the data more
 - More complex models
 - Tendency to overfit
 - Add more data points
 - “Unstable” models – model (and prediction) can change a lot

Model Complexity




$$\text{Expected Error} \approx \text{Bias} + \text{Variance}$$

Bias-Variance Tradeoff in Practice

- We saw that the classification error can be (informally) expressed in terms of bias and variance
- Reducing the bias and variance can reduce expected error!
- Different scenarios can lead to different actions for reducing the error
 - **High bias:** add more features
 - **High variance:** simplify the model, add more examples
- *How can we diagnose each one of these scenarios?*

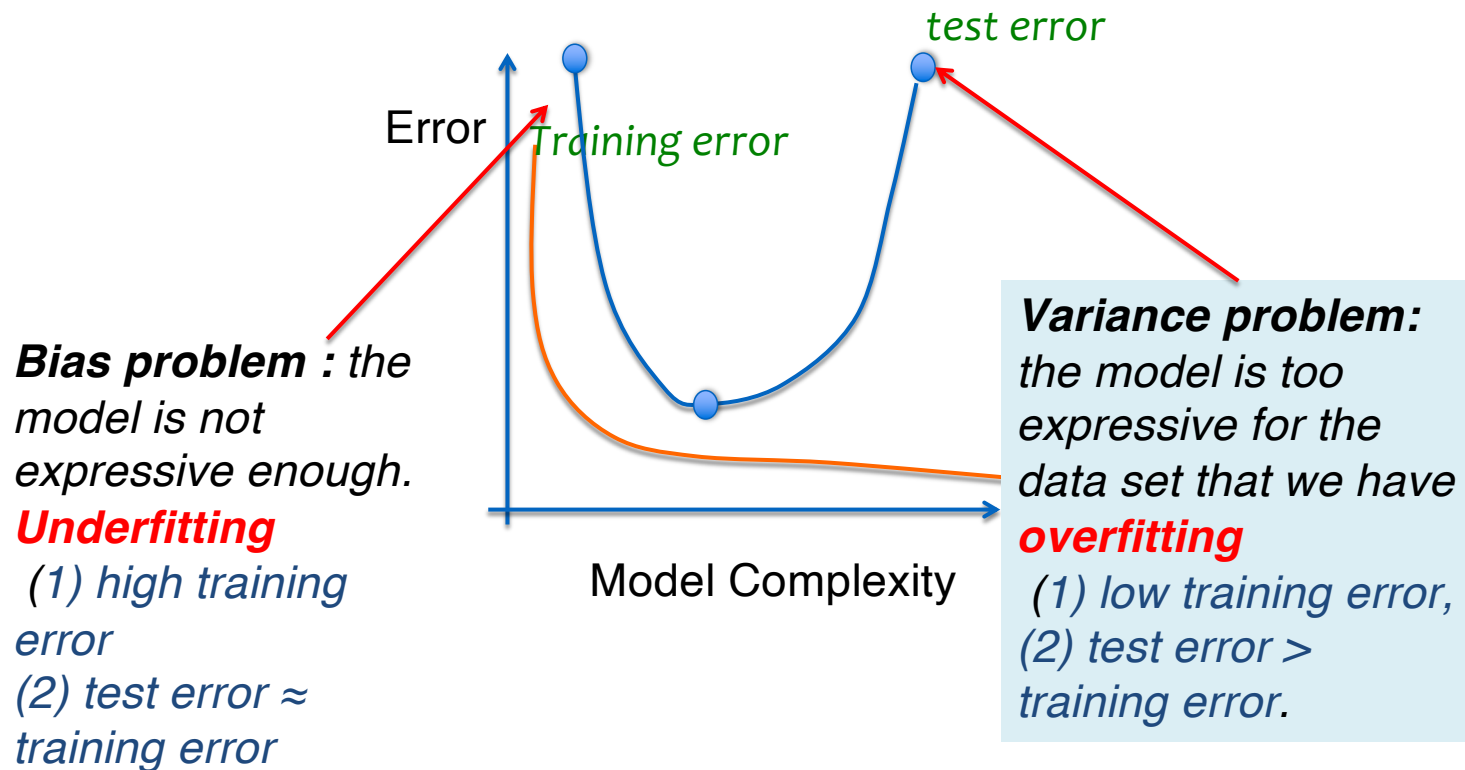
Overfitting




- Question 1: How do we detect overfitting?
- Question 2: How do we prevent / correct overfitting?

Bias-Variance Analysis

Interpolating over the points: two curves that we can use for **diagnosis**




Overfitting



- Question 1: How do we detect overfitting?
- **Question 2: How do we prevent / correct overfitting?**


Preventing / Correcting Overfitting



- Model Selection: Find a model that gives good test error
 - Often this is just tuning a “hyper-parameter” (e.g., k in k NN, bandwidth in kernel regression, C in Soft-SVM)
 - Sometimes more involved: e.g., post-pruning in decision trees
- A few ways to do this:
 - Validation set (or cross-validation)
 - Theoretical approaches [Structural Risk Minimization]

Training, Validation, Testing

- Split data set into **three** data sets: **training**, **validation**, **test**




Try different hyper-parameters
(for instance: $C=0.1$, $C=1$, $C=10$ for SVM, or $k=1,2,3,\dots,n$ for kNN)



Report test error rate for the hyper-parameter value
that gave smallest validation error rate

Cross-validation

- Problem: What if the training set is an “unlucky” distribution
 - Error on the test set doesn’t match real data
- Solution: Use *all* of the data as test data
 - But then we don’t have any data to train on!
- Instead, cross-validation
 - Multiple training/test runs
 - Each uses a different subset as test data



Cross-validation

- **K-fold cross-validation:**

- Randomly partition the training data into K equal-size subsets S_1, \dots, S_K

- For each i

- Train on $S_1 \cup \dots \cup S_{i-1} \cup S_{i+1} \cup \dots \cup S_K$ (all the data except S_i)
- Call this classifier \hat{h}_i
- Evaluate error rate on S_i : $\text{error}_{S_i}(\hat{h}_i)$

- Return the average: Cross-validation error $= \frac{1}{K} \sum_{i=1}^K \text{error}_{S_i}(\hat{h}_i)$


- We can use this to tune hyper-parameters

- For each setting of the hyper-parameters
- Run K-fold cross validation
- Choose the hyper-parameter values with lowest cross-validation error
- Retrain on the **entire** training set, using the chosen hyper-parameter values

Cross-validation

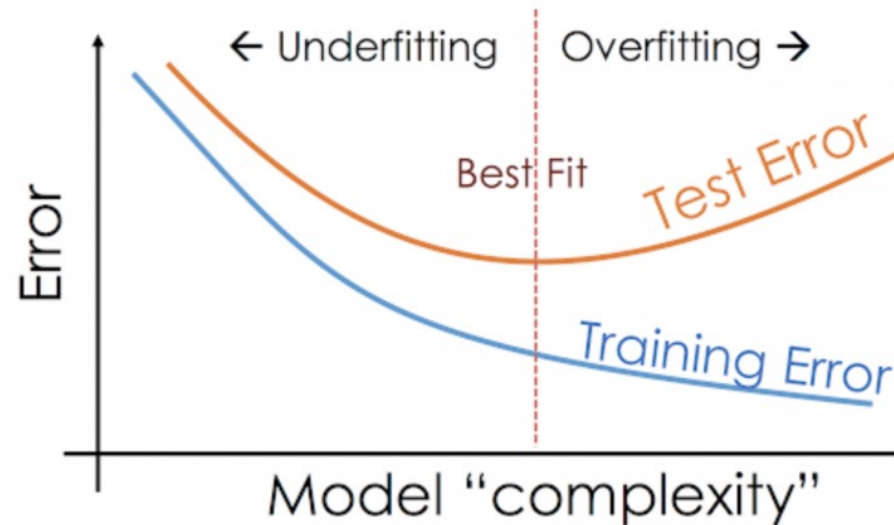
- How do we pick K ?
- Most popular in practice: $K=10$
- If $K=n$, it's called “leave one out” cross-validation:
 - Every training point (x_i, y_i) gets its own fold $S_i = \{(x_i, y_i)\}$
 - But this is computationally expensive
 - Also can sometimes have higher variance
- To estimate the error rate in the end, we would still need a separate held-out test set

Preventing / Correcting Overfitting



- Other approaches:
 - **Regularization** (e.g., minimizing $\|w\|$ in SVM optimization)
 - Dimensionality reduction / feature selection

Preventing / Correcting Overfitting



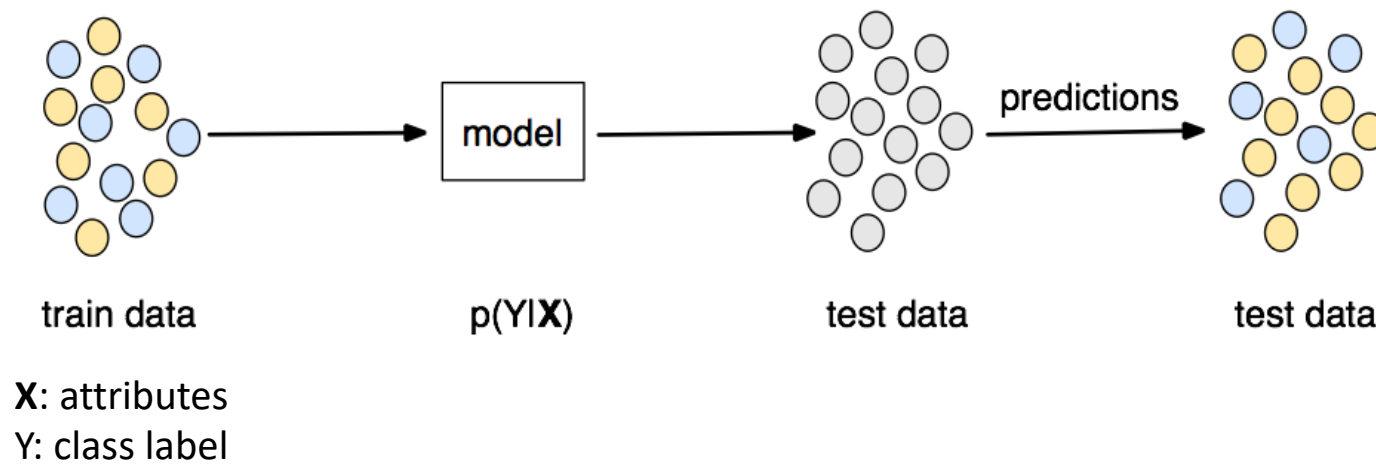
- Other approaches:
 - Regularization (e.g., minimizing $\|w\|$ in SVM optimization)
 - Dimensionality reduction / feature selection
- Note: sometimes the appropriate units for x-axis aren't easy to identify.
 - Sometimes large neural networks overfit less than smaller
 - Possibly because the optimization finds good solutions more easily

Ensemble methods

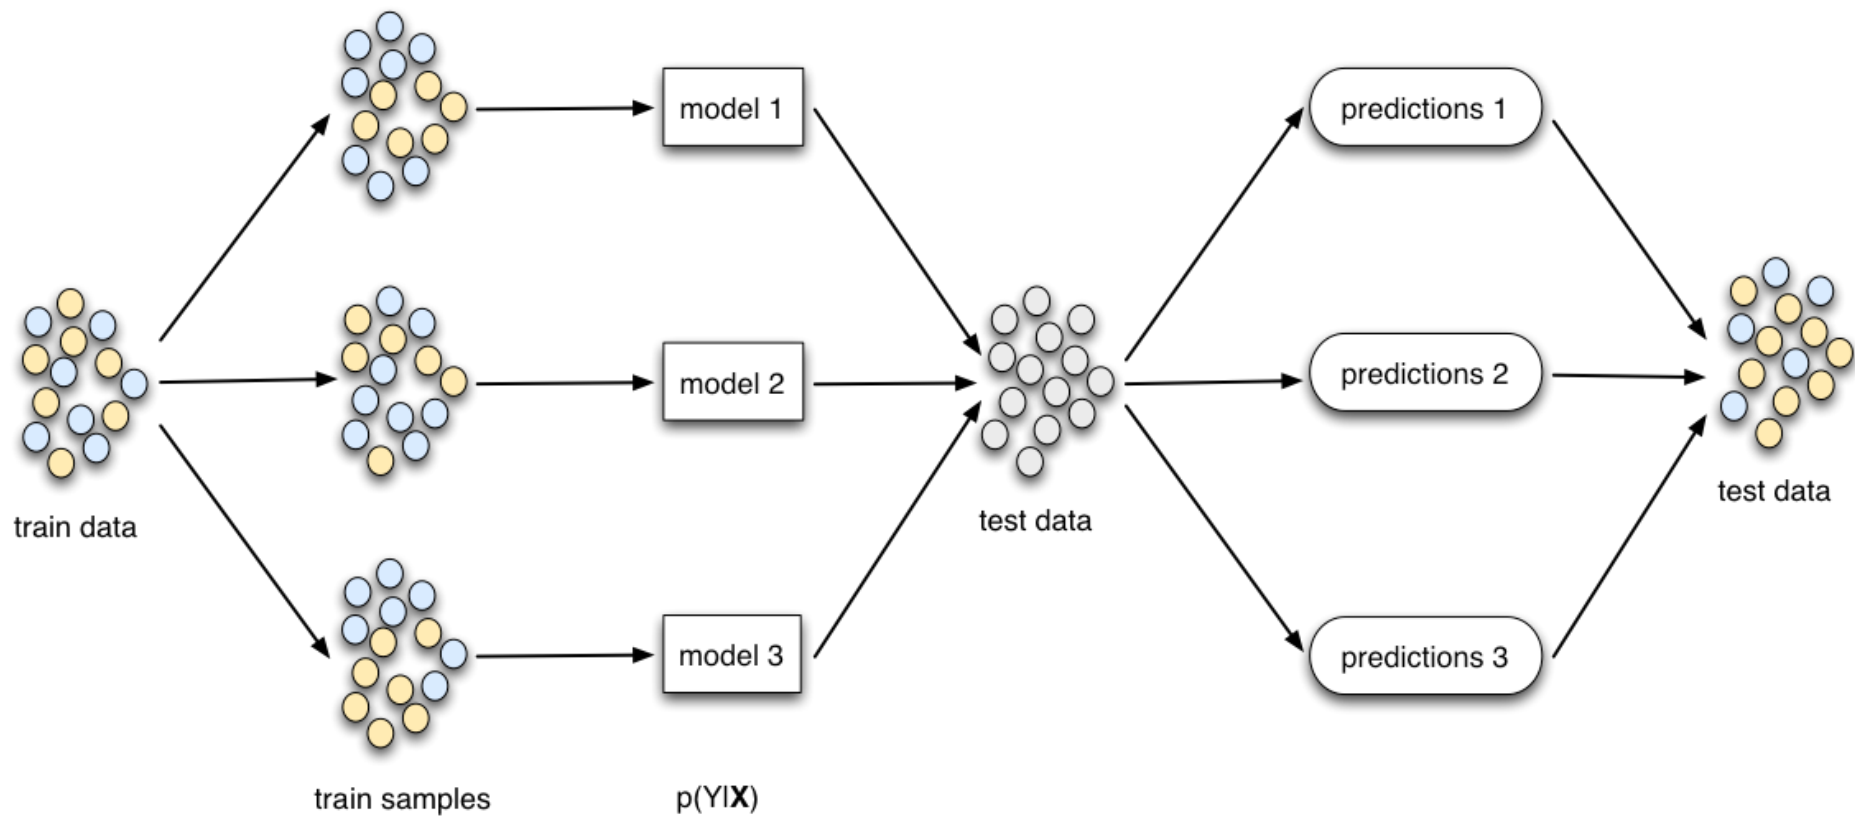
Ensemble methods

- Motivation
 - Too difficult to construct a single model that optimizes performance (why?)
- Approach
 - Construct many models on different versions of the training set and combine them during prediction
- Goal: reduce bias and/or variance

Conventional classification




Ensemble classification



Why not choose the best classifier?


- Best on what?

- Training data?
- Test data?
- After cross-validation?



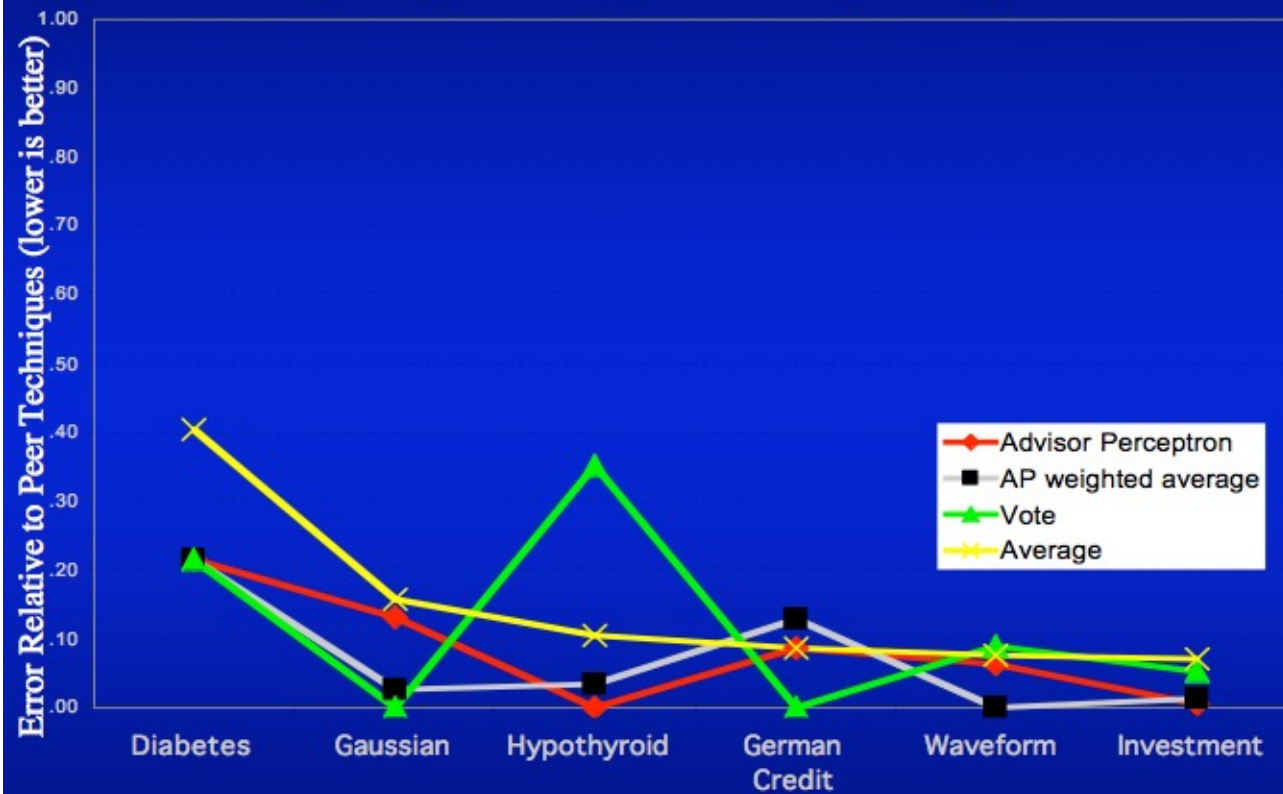
- *Think of as “multiple hypotheses, not sure which is best*

- Ensembles: use them all
- We can formally state this in terms of bias/variance reduction
 - Bias: Erroneous assumptions in the model
 - Variance: Sensitivity to small fluctuations in training data



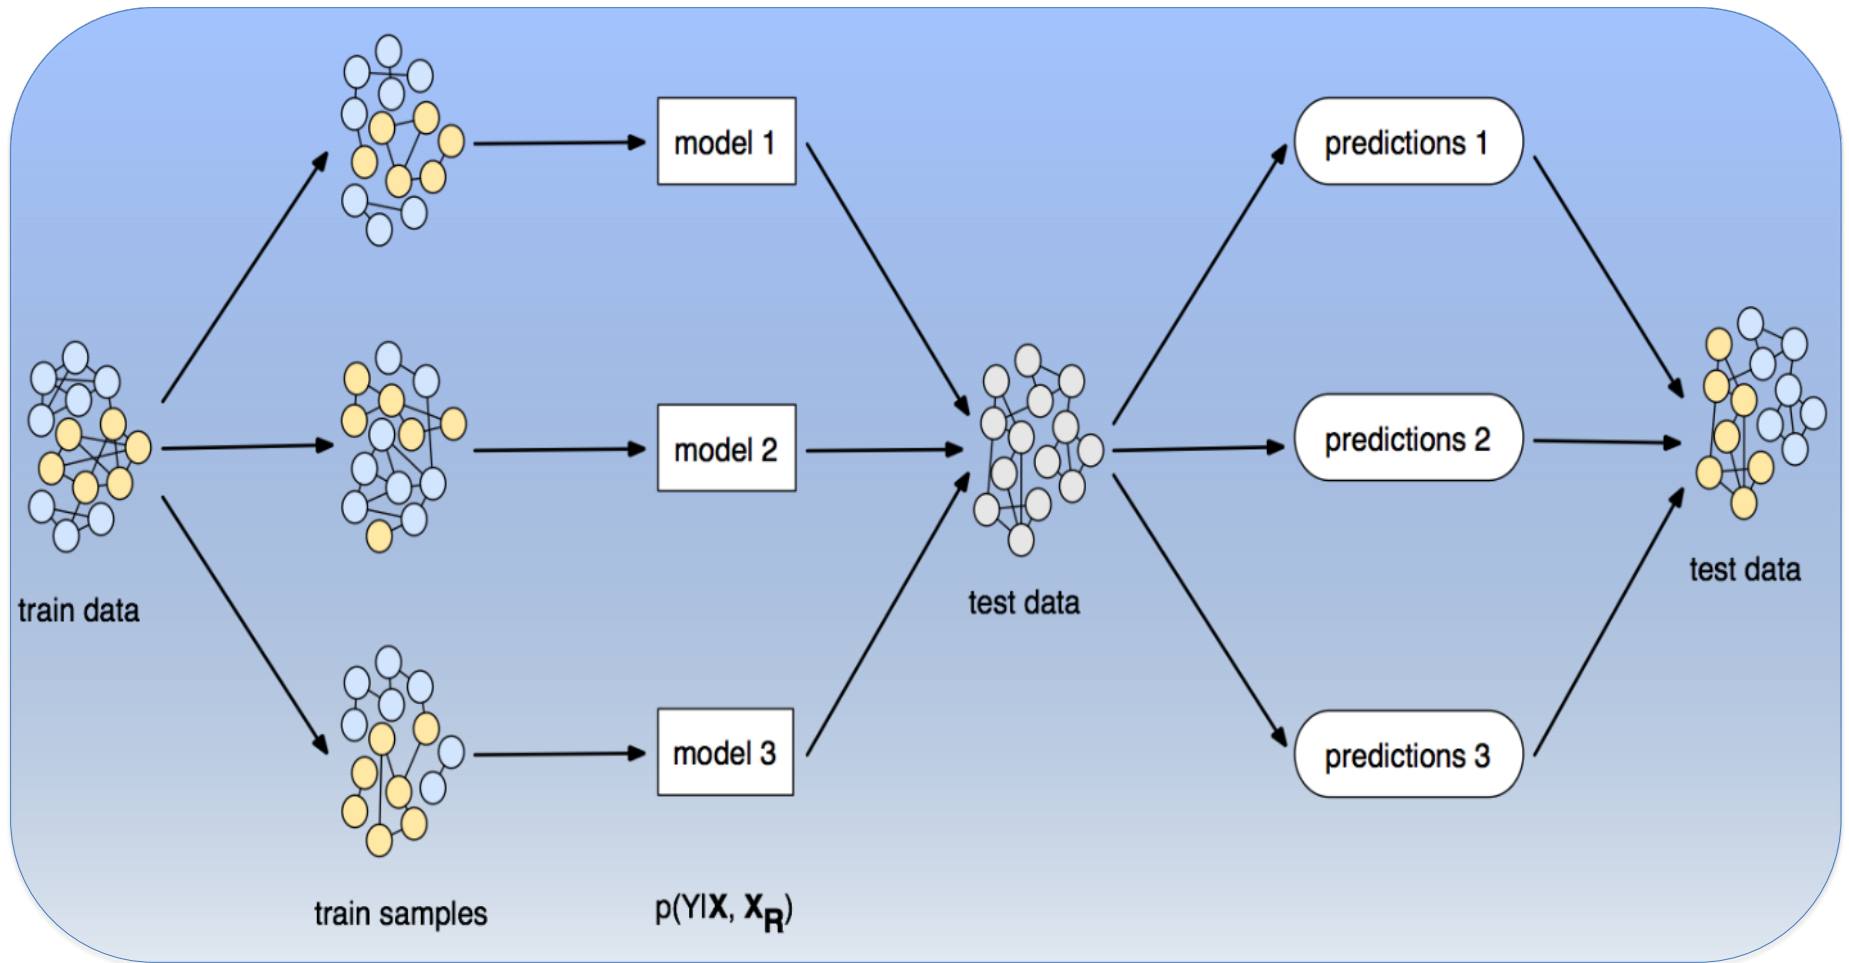
source: *Top Ten Data Mining Mistakes*,
John Edler, Edler Research)

Essentially every *ensemble* method improves performance

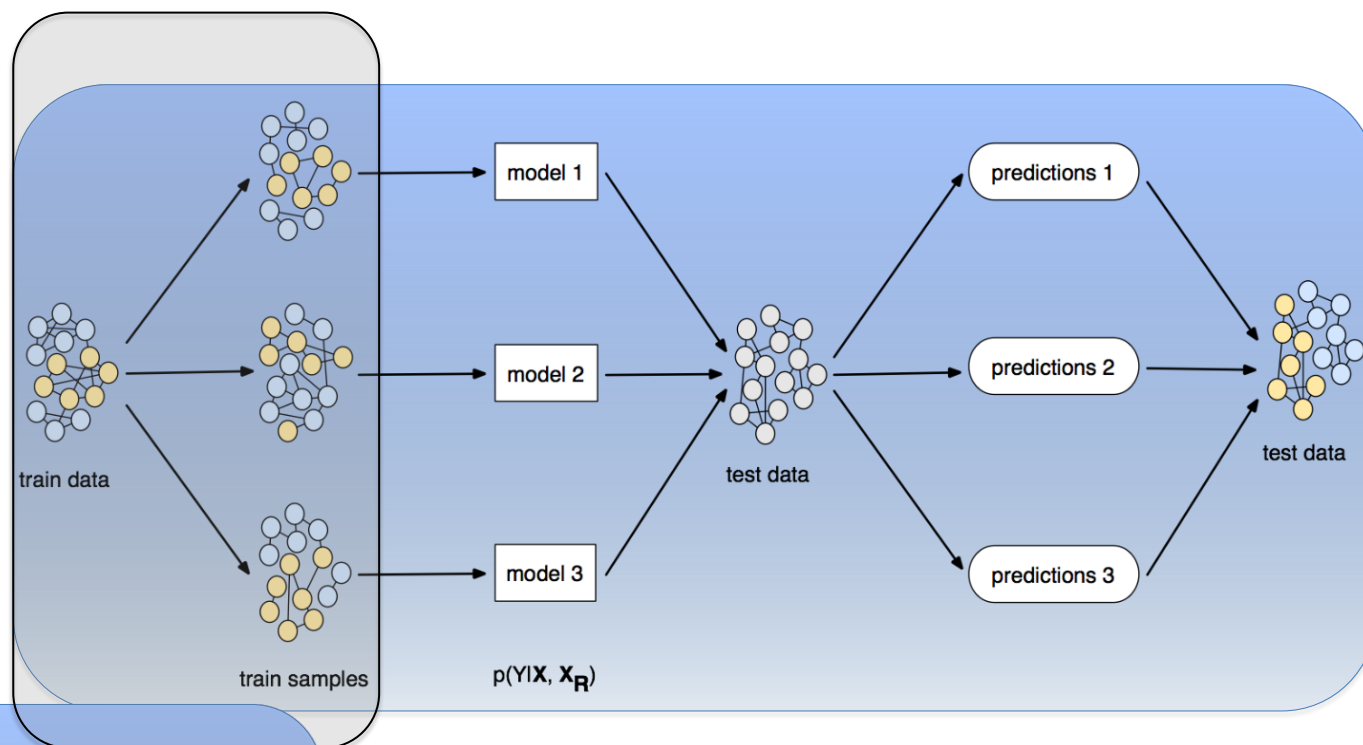


source: *Top Ten Data Mining Mistakes*,
John Edler, Edler Research)

Ensemble design



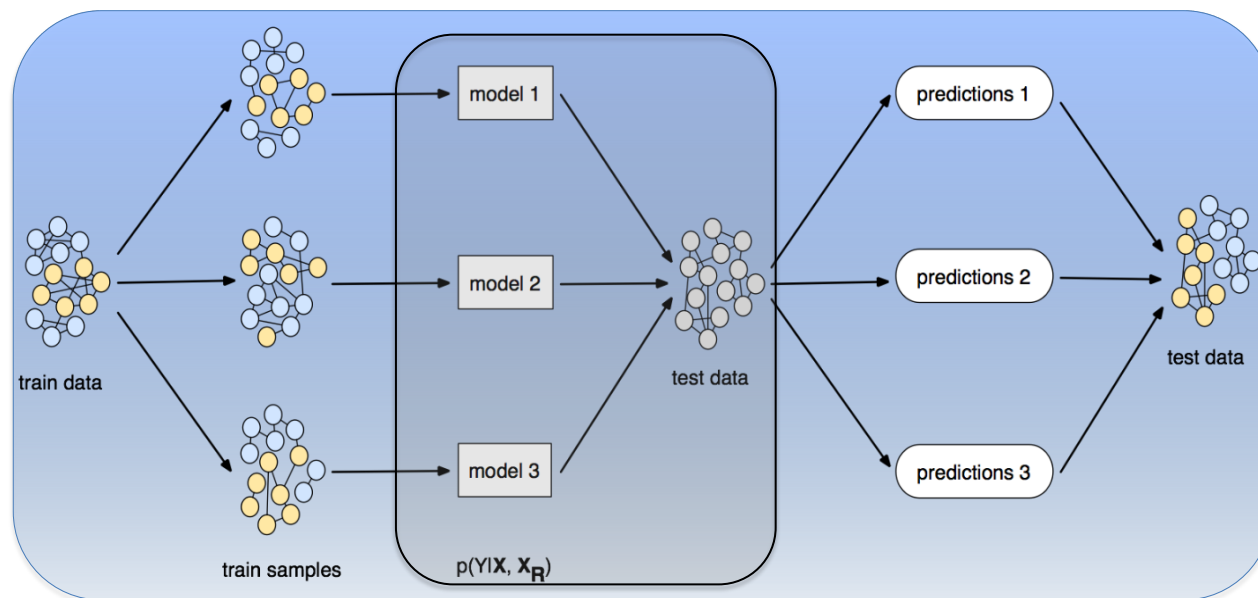
Ensemble design



TREATMENT OF INPUT DATA

- *sampling*
- *variable selection*

Ensemble design



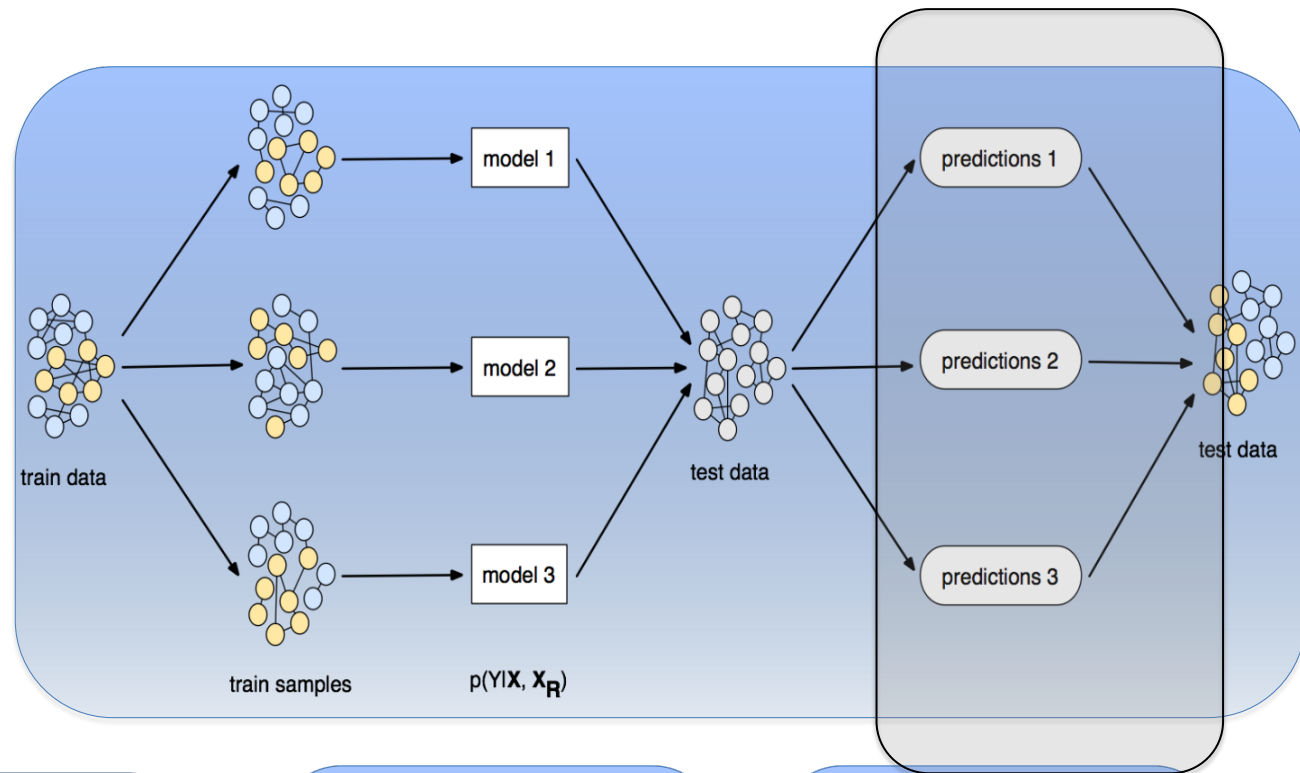
TREATMENT OF INPUT DATA

- *sampling*
- *variable selection*

CHOICE OF BASE CLASSIFIER

- *decision tree*
- *perceptron*
- ...

Ensemble design



TREATMENT OF INPUT DATA

- *sampling*
- *variable selection*

CHOICE OF BASE CLASSIFIER

- *decision tree*
- *perceptron*
- ...

PREDICTION AGGREGATION

- *averaging*
- *weighted vote*
- ...

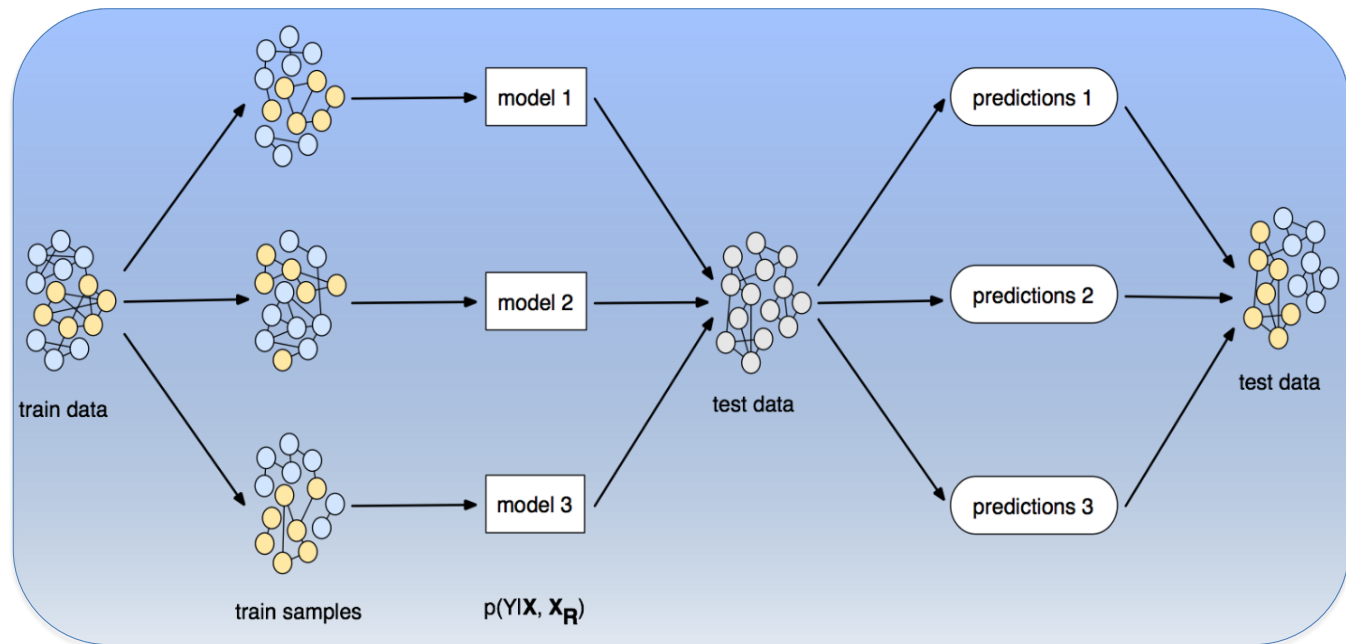
Bootstrap Sampling

- Generate multiple training sets from original training data
 - Sound familiar?
 - *Cross-validation*
- Key distinction: Sample *with replacement*

Bagging

- **B**ootstrap **agg**regating
- Main assumption
 - Combining many *unstable* predictors in an ensemble produces a *stable* predictor (i.e., reduces variance)
 - Unstable predictor: small changes in training data produces large changes in the model (e.g., trees)
- Model space: non-parametric, can model any function if an appropriate base model is used

Bagging



TREATMENT OF
INPUT DATA

- *sample with replacement*

CHOICE OF BASE
CLASSIFIER

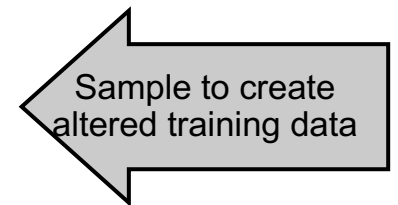
- *unstable predictor
e.g., decision tree*

PREDICTION
AGGREGATION

- *averaging*

Bagging

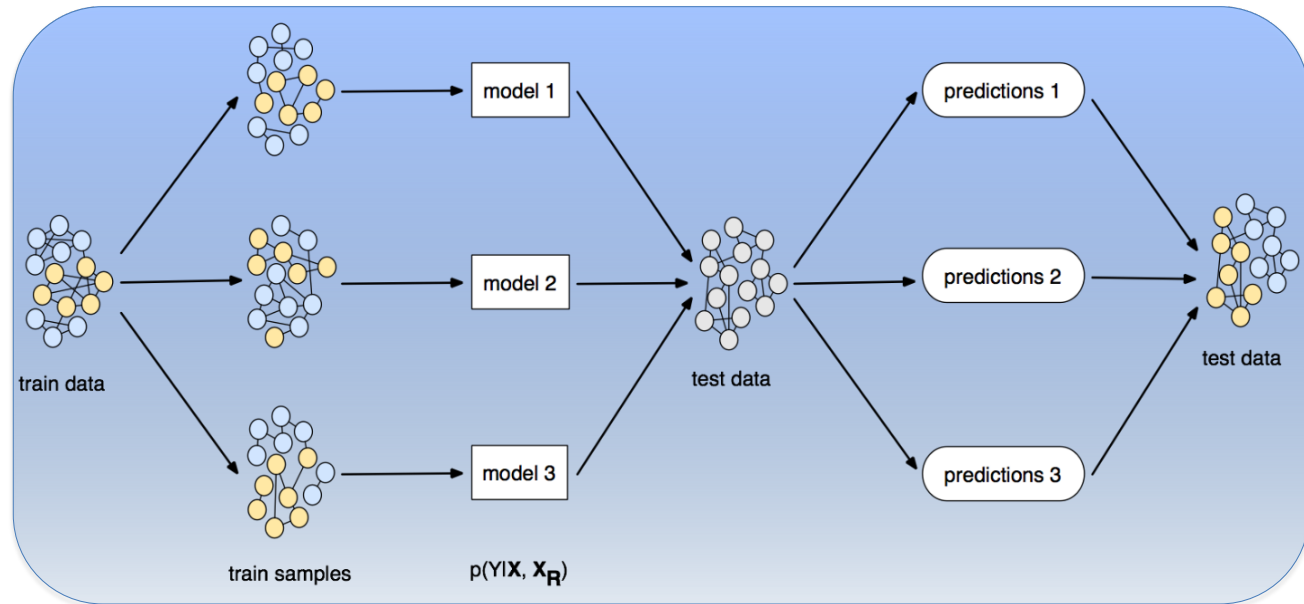
- Given a training data set $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$
- For $m = 1:M$
 - Obtain a bootstrap sample D_m by drawing N instances ***with replacement*** from D
 - Learn model M_m from D_m
- To classify test instance t , apply each model M_m to t and use majority prediction or average prediction
- Models have uncorrelated errors due to difference in training sets (each bootstrap sample has ~68% of D)



Boosting

- Main assumption
 - Combining many *weak* (but stable) predictors in an ensemble produces a *strong* predictor (i.e., reduces bias)
 - Weak predictor: only weakly predicts correct class of instances (e.g., tree stumps, 1-R)
- Model space: non-parametric, can model any function if an appropriate base model is used

Boosting



TREATMENT OF INPUT DATA

- *reweight
examples*

CHOICE OF BASE CLASSIFIER


- *weak predictor
e.g., decision
stump*

PREDICTION AGGREGATION

- *weighted vote*

Boosting: Adaboost

- Assign every example in D an equal weight ($1/N$)
- For $m=1:M$
 - Learn model M_m with D_m
 - Calculate the error of M_m and up-weight the examples that are incorrectly classified to form D_{m+1}
 - Normalize weights in D_{m+1} to sum to 1
 - Set $\alpha_m = \log((1-\text{err}_m)/\text{err}_m)$
- To classify test instance t , apply each model M_m to t and take weighted vote of predictions (ie. using α_m)



Adaboost Algorithm

- Takes training data (x_i, y_i) (y -1 or 1), weights w_i
 - Initialize weights to $1/n$
- For $m=1..M$
 - Learn classifier f_m
 - $Error_m = \sum_{i=1}^n w_i^m \mathbf{I}\{f_m(x_i) \neq y_i\}$
 - Compute classifier coefficient $\alpha_m = \frac{1}{2} \log \frac{1 - Error_m}{Error_m}$
 - Update weights $w_i^{m+1} = \frac{w_i^m \exp(-\alpha_m y_i f_m(x_i))}{\sum_{j=1}^n w_j^m \exp(-\alpha_m y_j f_m(x_j))}$
- Final classifier $f^*(x) = \text{sign}(\sum_{m=1}^M \alpha_m f_m(x))$

Boosting Caveats

- While theoretically sound, Adaboost not that robust to noisy labels
 - Weights of mislabeled data grow until classifier fits the noise
- **Must** use *weak* classifiers
 - Otherwise easily overfits training data

Random Forests

- Problem: Decision Trees prone to overfitting
- Solution: Decision tree on fewer features
- Ensemble idea
 - Randomly select subsets of features
 - Choose best candidate split from just within subset
- Algorithm the same as standard decision tree, except instead of applying information gain / gini index / ..., first randomly select subset, then apply
 - All features (except the one use) passed to the next level

Ensemble summary

- Two approaches for Ensemble learning:
 - Boosting – reduce bias
 - Bagging – reduce variance
- Applicable in different situations