

# Data Mining & Machine Learning

---

CS37300  
Purdue University

Oct 30, 2023

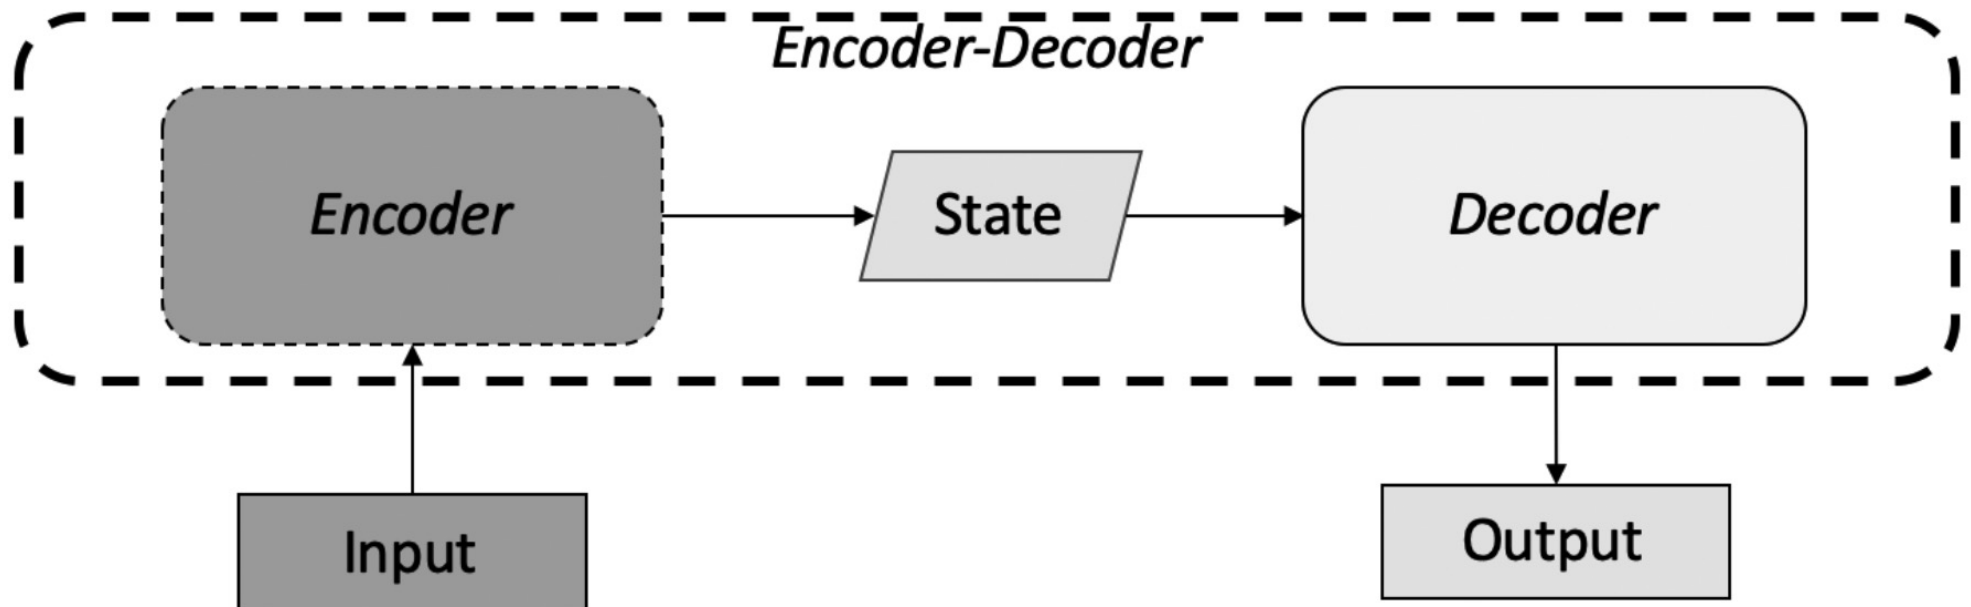
# Transformers

---

# Sequence-to-sequence models

---

- Encoder-decoder architectures

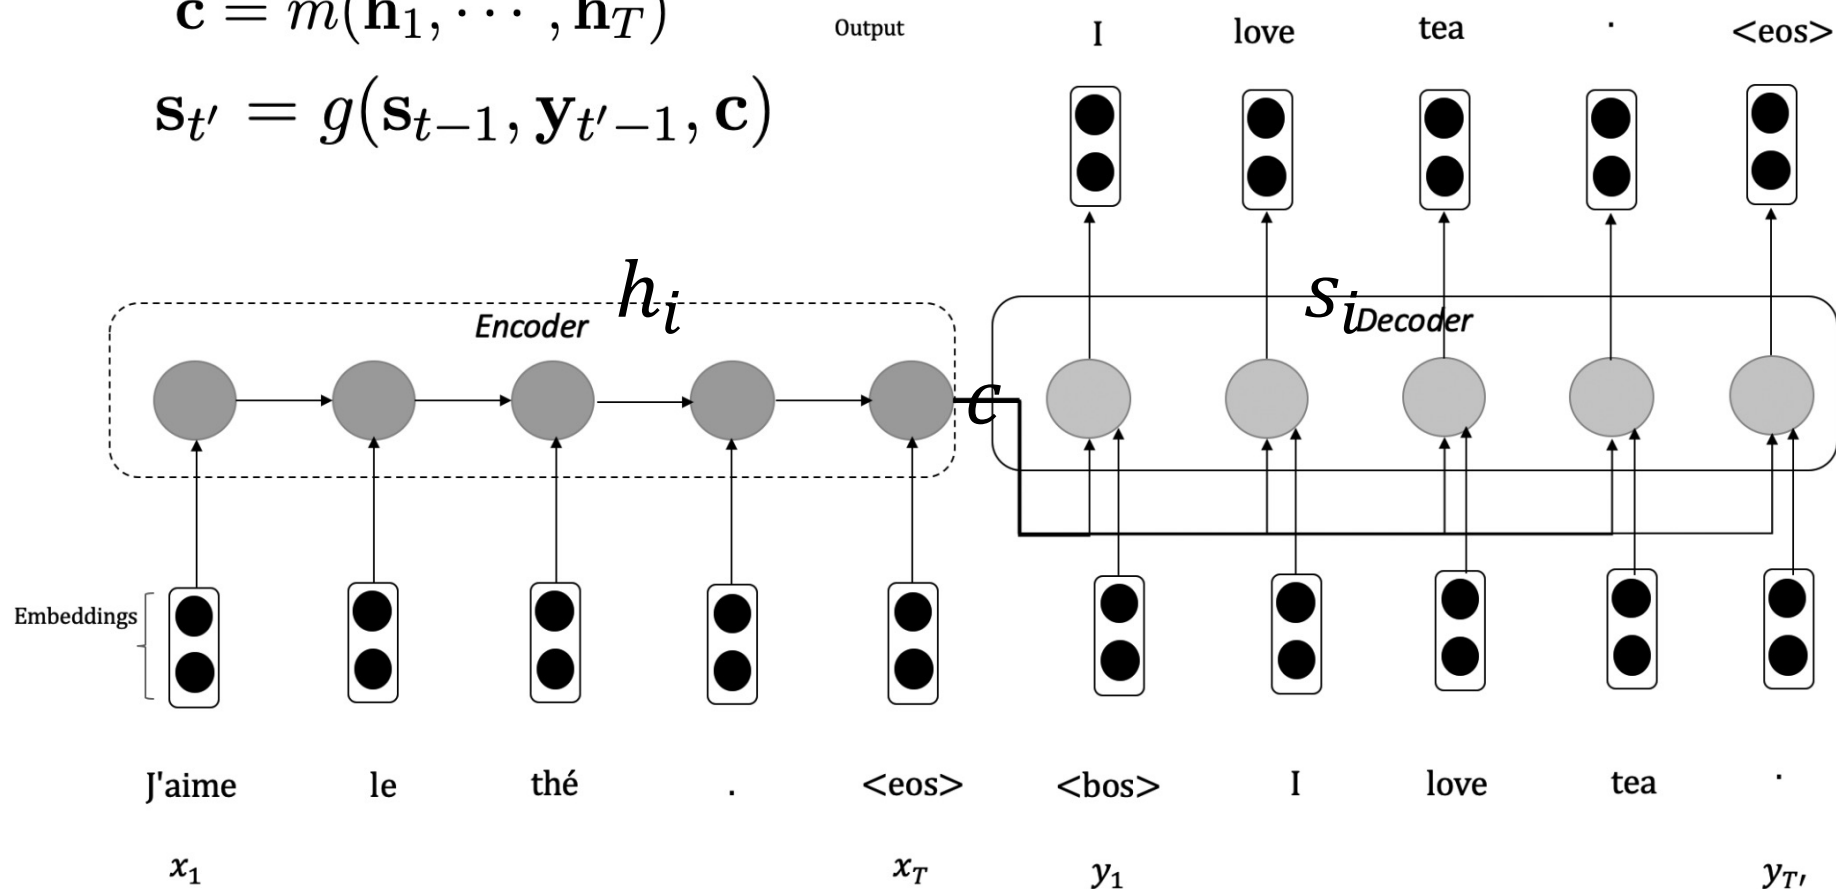


# Combine two RNNs to build seq2seq models?

$$\mathbf{h}_t = f(\mathbf{h}_{t-1}, \mathbf{x}_t)$$

$$\mathbf{c} = m(\mathbf{h}_1, \dots, \mathbf{h}_T)$$

$$\mathbf{s}_{t'} = g(\mathbf{s}_{t-1}, \mathbf{y}_{t'-1}, \mathbf{c})$$




# Limitations of RNNs for seq2seq

---

- A single context vector linking two RNNs may compress too much (loss of information)
- Could potentially link hidden vars of encoder to hidden vars of decoder
- Generally for RNNs:
  - Exploding/vanishing gradients
  - Recurrence relationships make it difficult to parallelize

# Attention mechanism

- Focus on selective information



# Attention mechanism contd

---

- Memory: Stores key-value pairs  $(k_i, v_i)$
- For a new query  $q$ , calculate its similarity with stores keys:
  - $b_i = \text{sim}(q, k_i)$ . Normalize so that  $\mathbf{b}$  is a probability distribution
- Based on similarity, form a “weighted value output”
  - $o = \sum_i b_i v_i$

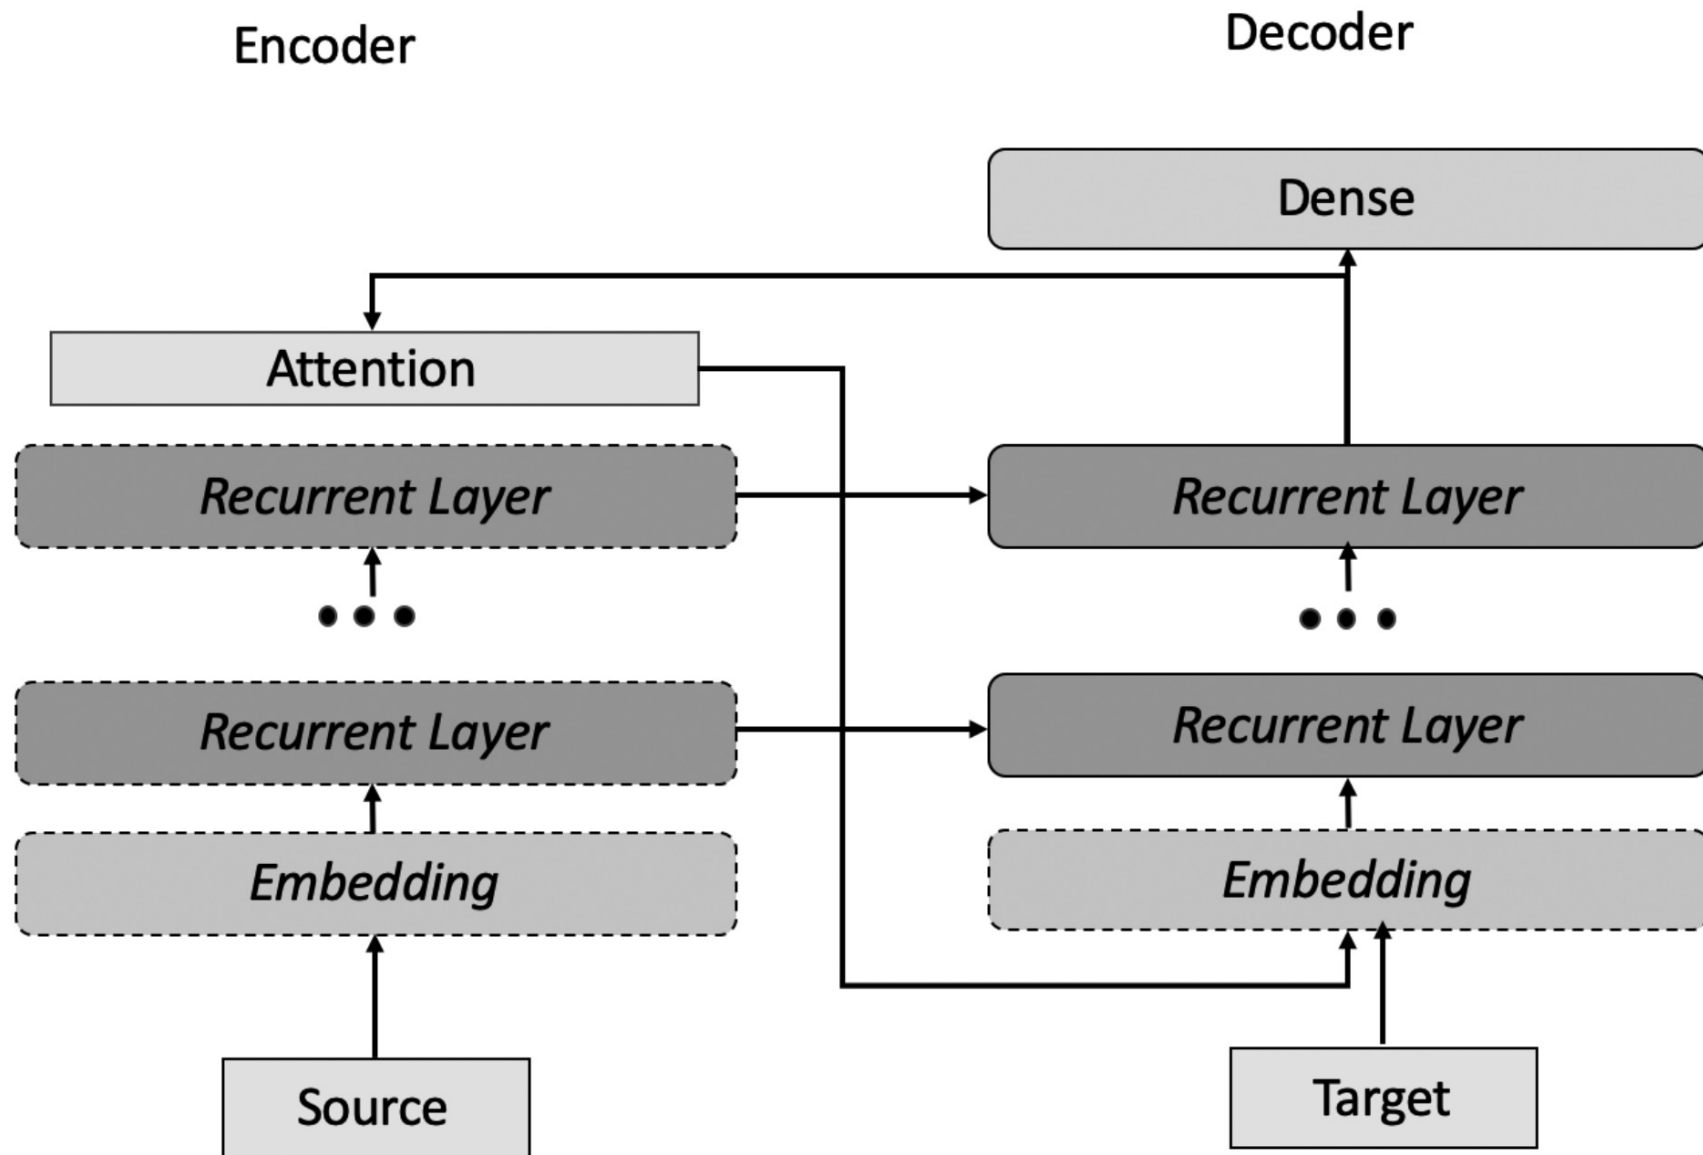
# Similarity function

---

- Used for calculating similarity in score and query
- Any suitable kernel function
  - Dot product
  - Mahalanobis similarity
  - Latent variable based



# Encoder-decoder with attention



# Implementation of attention

---

- Several Recurrent layers
- keys, values extracted from encoder states
- Query extracted from decoder's state at time/step  $t-1$
- Context output from attention used for next state  $t$  of the decoder
- Context variable here is a weighted average from several keys/values extracted from the encoder states

# Transformers

---

- Combine CNNs (parallelize computations) and RNNs (capture sequences)
- Main components:
  - Inputs/Outputs tokenized and positionally embedded
    - Orders could be handled in RNNs, but here the feature vectors encode the relative positions
  - Masked multi-head attention
  - [contd later]

# Self-attention

---

- Model dependencies amongst in the input embeddings
- Obtain key,value,query triplets from the same input embedding  $x_i$  by learning projection matrices  $W_k, W_v, W_q$
- The query vector  $q_i = W_q x_i$ , is combined with all other key vectors in the input to generate the output for its own self  $\sum q_i k_i^T$
- The outputs are then generated by combining normalized weighted self-query with the value, where  $d_k$  is the size of the key vector

$$attention(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left( \frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V}$$

# Multi-head attention

---

- Combine several individual attentions learnt above
- Analogous to different convolution filters in CNNs

$$head_i = attention(\mathbf{W}_q^i \mathbf{Q}, \mathbf{W}_k^i \mathbf{K}, \mathbf{W}_v^i \mathbf{V})$$

$$multihead(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \mathbf{W}_O \text{concat}(head_1, \dots, head_h)$$

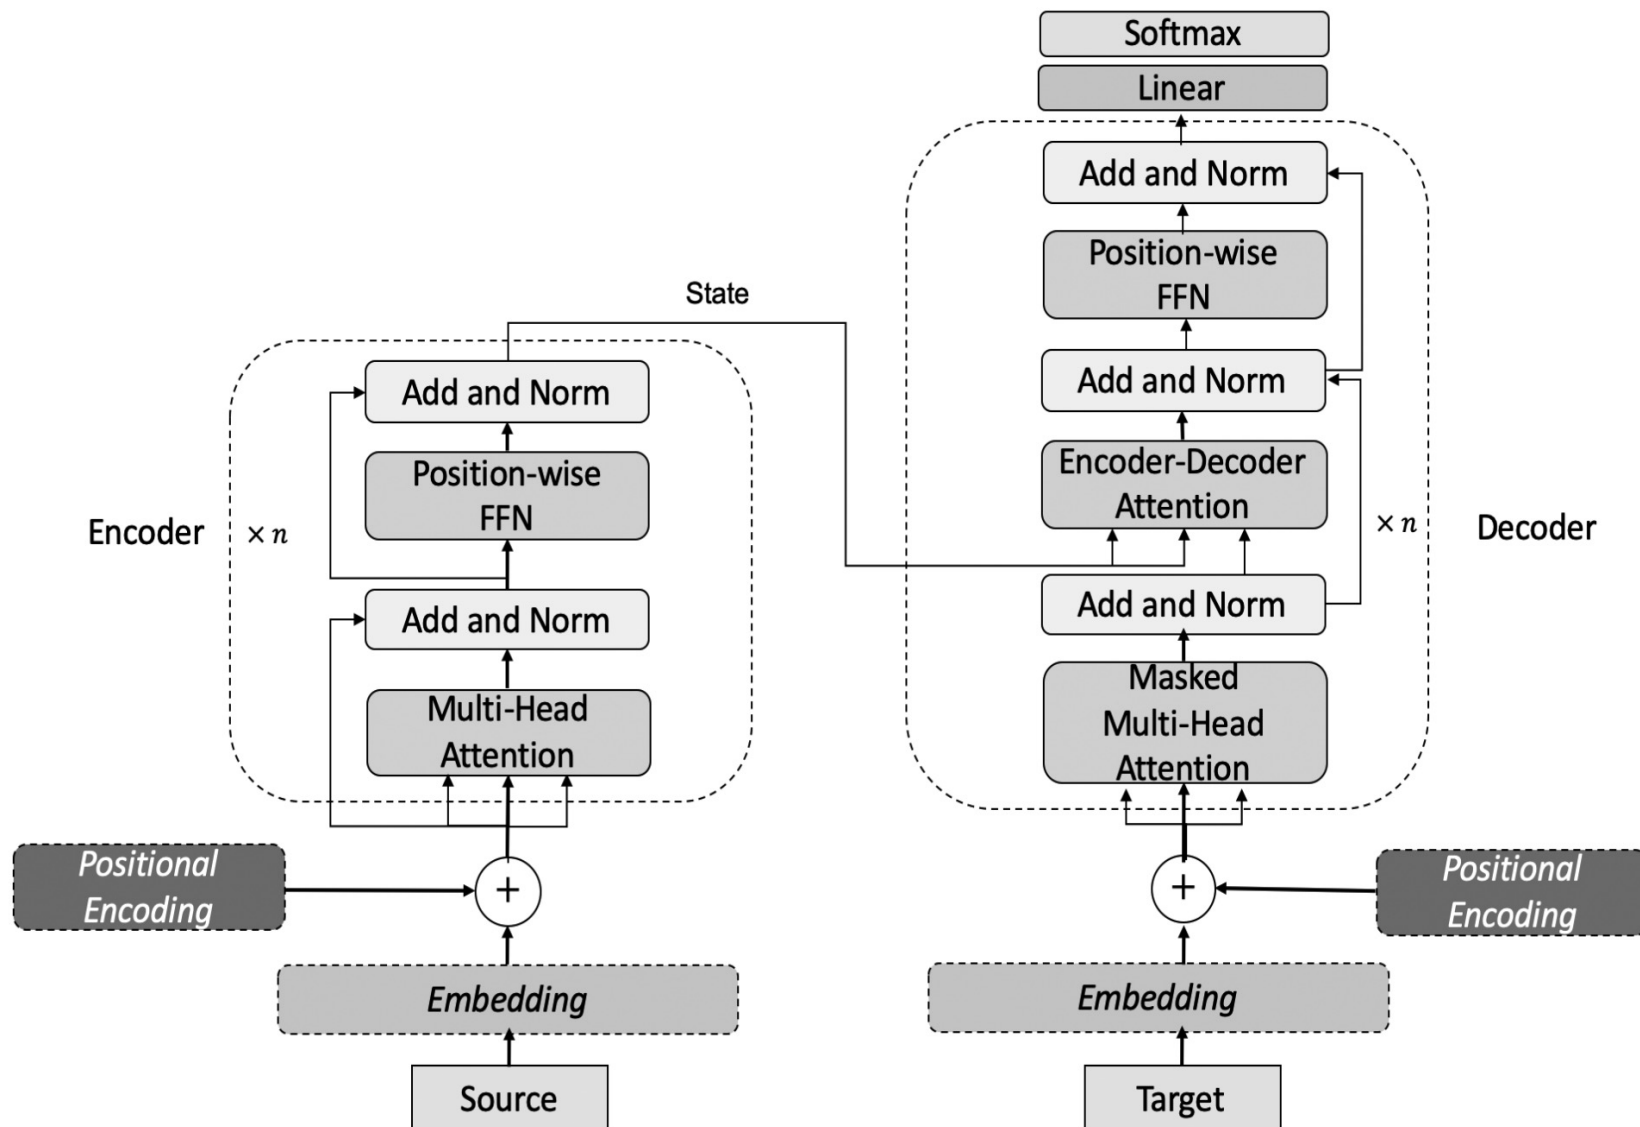
- Masked attention to remove the influence of future tokens for the current token in learning
- Encoder only sees attention learning from inputs
- Decoder sees query vectors from target outputs as well output from the encoders

# Transformer components (contd)

---

- Inputs/Outputs embeddings
- Masked multi-head attention
- Feed-forward networks after the attention mechanisms
- Residual (skip) connections
- Layer normalization – for faster convergence

# Components of a transformer block



# Optimization of NNs

---

- An accelerated variant of Batch SGD
  - Can parallelize gradient evaluations across multi-core architectures
  - Batch size has a regularizing effect
- Adaptive learning rate
- Challenges
  - Local minima, saddle points
  - Vanishing/Exploding gradients
  - Ill-conditioned Hessians